

Correct Zipcodes using Serverless Golang

Max and Sebastian

Nov 08 2019

thinkproject

The bottom right corner of the slide features several overlapping geometric shapes. There is a white parallelogram, a teal parallelogram, and a green parallelogram, all slanted at an angle, creating a modern, abstract design.

What the heck

Share our passion about Go based Microservices:

- ▶ implementing a REST service
- ▶ ... in go
- ▶ ... using a Google Cloud Function

Show me what you've got

use for example:

```
curl -X POST ... -d '{"zipCode":"72205", "placeName":"Barlin"}'
```

to get:

```
"distance":2,  
"percentage":81,  
"place":{  
  "countryCode":"DE",  
  "zipCode":"12205",  
  "place":"Berlin",  
  ...  
  "latitude":"52.434",  
  "longitude":"13.2945"  
}
```

What's the point?

- ▶ What's the fuss about FaaS?
- ▶ Support EPLASS address detection.
- ▶ Let's play go.

TOC

1. FaaS
2. V0, V1 and V2 of a slack command
3. GCF in Production
4. Wrapup

TOC

1. **FaaS**
2. V0, V1 and V2 of a slack command
3. GCF in Production
4. Wrapup

FaaS




Figure 1: *aaS Stacks

FaaS

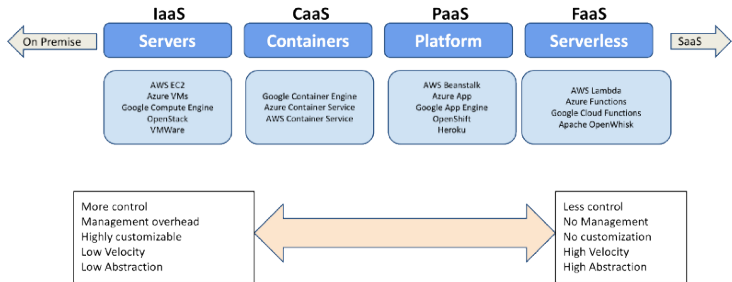


Figure 2: *aaS Provider

FaaS here

- ▶ AWS Lambda, **Google Cloud Functions**, MS Azure Functions, . . .
- ▶ Node.js, Python, **Go 1.11.5**
- ▶ Google Cloud Platform, **Google Cloud SDK**, (mirrored) Source Control

TOC

1. FaaS
2. **V0, V1 and V2 of the zipchecker**
3. GCF in Production
4. Wrapup

V0 Hello World

v0 implements a service that responds with “Hello World”:

- ▶ project layout
- ▶ implementation
- ▶ building and running

V1 Business Logic

v1 extends v0 by zipchecker business logic:

- ▶ implementation
 - ▶ request processing
 - ▶ marshaling and unmarshaling
 - ▶ embedding statics
 - ▶ constructors
 - ▶ Levenshtein distance
- ▶ testing

V2 GCP Deployment

v2 extends v1 by GCP deployment:

- ▶ GCP preparation
- ▶ GCP admin console
- ▶ GCP deployment
- ▶ GCP logs

TOC

1. FaaS
2. V0, V1 and V2 of a slack command
3. **GCF in Production**
4. Wrapup

GCF in Production – Scaling

- ▶ scales by creating new function instances
- ▶ the total number of function instances can be limited
- ▶ function instances are reused
- ▶ global scope may be used to cache across function invocations
- ▶ concurrent requests are processed by different instances
- ▶ response-time depends on hot- or cold start

GCF in Production – Pricing

- ▶ \$0.40 / 10^6 invocations
- ▶ \$0.0000025 / GB-Second memory
- ▶ \$0.0000100 / GHz-Second CPU
- ▶ \$0.12 / GB Outbound Data (Egress) traffic

“Free Tier” per month:

- ▶ 2 million invocations
- ▶ 400,000 GB-seconds
- ▶ 200,000 GHz-seconds
- ▶ 5 GB Egress traffic

see: [Cloud Functions Pricing](#)

GCF in Production – Price Example

based on 2ms (i.e. 100ms) + 1KB traffic at 128MB and 200MHz:

- ▶ CPU: $\frac{2 \cdot 10^8 \text{ MHz s}}{0.1 \text{ s} \cdot 200 \text{ MHz}} = 10 \cdot 10^6$
- ▶ memory: $\frac{400,000 \cdot 1024 \text{ MB s}}{0.1 \text{ s} \cdot 128 \text{ MB}} = 32 \cdot 10^6$
- ▶ traffic: $\frac{5 \cdot 1024 \cdot 1024 \text{ KB}}{1 \text{ KB}} \approx 5.2 \cdot 10^6$
- ▶ invocations: $2 \cdot 10^6$

$\min(10, 32, 5.2, 2) \rightarrow 2 \cdot 10^6$ free invocations $\equiv \sim \$4.88$

TOC

1. FaaS
2. V0, V1 and V2 of a slack command
3. GCF in Production
4. **Wrapup**

Wrapup – Things to do

- ▶ trees of functions
- ▶ provided services

Wrapup – The next big thing?

- ▶ interesting idea
- ▶ for small services: easy implementation and deployment

however:

- ▶ implementation becomes even more fragmented
- ▶ “overly distributed”
- ▶ difficult to test
- ▶ high delay for logs (sometimes 5-10s)
- ▶ cost: hard to predict and might get out of hand

Wrapup – Readings

- ▶ Google Cloud Functions Tutorial Series
- ▶ this talk, the code, etc.
- ▶ Service Setup with Gin, Auth0
- ▶ Service Monitoring with Go and DataDog

end